

Akash Ghadage

Senior Data Engineer

✉ mailtoakashghadage@gmail.com |  LinkedIn: Akash | +91-9975796681 | 📍 Pune, India

SUMMARY

- Databricks Certified Azure Data Engineer with **4+ years of experience** designing, building, and optimizing scalable **batch and streaming ETL/ELT pipelines** on **Azure and Databricks**. Specialized in **Lakehouse architecture** and large-scale data processing, building reliable data platforms for **enterprise analytics, reporting, and AI/ML workloads** across **Fortune 500 clients**.
- Reduced pipeline execution time by **30%** through **Apache Spark optimization** and cut query latency by **40%** through **Delta Lake optimization techniques**. Skilled in **Azure Databricks, Apache Spark, PySpark, Python, SQL, Azure Data Factory, and Delta Lake**, with experience building **event-driven architectures** and delivering **cost-efficient, scalable data solutions**.

TECHNICAL SKILLS

- **Programming:** Python, PySpark, SQL, Spark SQL
- **Big Data & Streaming:** Databricks, dbt, Delta Lake, Apache Spark, Unity Catalog, Apache Kafka, Spark Structured Streaming, Databricks Auto Loader, Performance Tuning & Optimization
- **Azure Services:** Azure Data Factory, Azure Key Vault, Azure SQL, Azure Cosmos DB, ADLS Gen2
- **Data Architecture:** Lakehouse Architecture, Medallion Architecture, Data Governance
- **APIs & Integration:** REST APIs
- **CI/CD / Agile:** Azure DevOps, Git, GitHub, Jira
- **Generative AI:** LLM Fundamentals, Prompt Engineering, Retrieval-Augmented Generation (RAG), GitHub Copilot, Databricks Genie, Databricks Mosaic AI (Vector Search, Agent Framework)

CERTIFICATIONS

- [Databricks Certified Data Engineer Professional](#)
- [Databricks Certified Generative AI Engineer Associate](#)
- [Microsoft Certified: Azure Data Engineer Associate \(DP-203\)](#)
- [Microsoft Certified: Azure Fundamentals \(AZ-900\)](#)
- [Microsoft Certified: Azure AI Fundamentals \(AI-900\)](#)
- [Microsoft Certified: Azure Data Fundamentals \(DP-900\)](#)

EXPERIENCE

Capgemini | Azure Data Engineer | Pune, India

Mar 2025 - Present

Enterprise Data.One Lakehouse (Bayer):

- Contributed to the development and operations of a data lakehouse built on Azure Databricks to process ERP and non-ERP data, supporting AI/ML and BI use cases.
- Developed a robust, **metadata-driven ETL framework** leveraging **Medallion Architecture** to process 17K+ tables and terabytes of data using Databricks Auto Loader, Delta Lake, Python, PySpark, and Apache Spark
- Implemented **2 key Spark optimization strategies, fine-tuned Spark jobs**, resulting in a **30-40% reduction in operational costs** and enhanced performance.
- Transitioned 200+ Databricks workloads from **Databricks Runtime (DBR) 13.3 LTS to 15.5 LTS**, improving performance and stability (**Awarded by TCS**).
- **Implemented Delta Lake optimization techniques**, including partitioning, data skipping, OPTIMIZE, Z-Ordering, and liquid clustering, **reducing query execution time by 40%**.
- Created Python automation to monitor Databricks Personal Access Token (PAT) expiry and generate alerts, **reducing authentication and integration failures by 100%** across connected systems.
- Designed a PySpark-based data validation and recovery framework that ensures data quality, integrity, and reliable downstream reporting.

Customer 360 – CDIS Event-Driven Integration Platform:

- Developed an event-driven customer master data integration platform using **Azure Databricks, Apache Kafka, Spark Structured Streaming, GraphQL, and REST APIs** to enable bidirectional synchronization between **Customer 360 (C360)** and **CDIS**.
- Engineered scalable **PySpark** pipelines using **Spark Structured Streaming** and **Delta Lake** to process **40K Kafka events** and **50K customer records** every **30 minutes**, ensuring reliable, fault-tolerant bidirectional data synchronization.
- Implemented automated configuration synchronization, retry mechanisms, and custom reprocessing workflows using **Python** and **Delta Lake**, eliminating manual intervention and improving pipeline reliability.

Tata Consultancy Services | Azure Data Engineer | Pune, India

Aug 2022 - Mar 2025

C360 Lakehouse Azure-to-Azure Migration (J&J → Kenvue)

- Migrated **Azure Databricks** and **Azure Data Factory** workloads to a shared Azure platform, **improving processing efficiency by 30%**.
- Built scalable **Azure Data Factory** and **Databricks** pipelines using **Python, PySpark, and Spark SQL** to ingest and transform **1 TB of enterprise data daily** into **ADLS Gen2** across **L0 (Raw), L1 (Harmonized), and L1+ (Semantic)** layers, **achieving 99.9% data availability**.
- Used **Azure DevOps** and **Git** for version control, CI/CD, deployment automation, and release management of **Azure Databricks Jobs** and **Azure Data Factory** pipelines.

EDUCATION

B.Tech in Computer Science and Engineering — CGPA: 9.14/10

Sanjeevan Engineering and Technology Institute

July 2018 – July 2022

Kolhapur, Maharashtra

ACHIEVEMENTS

- **Xcelerate Warrior Certificate** - TCS
- **2X Star Team Award** for exceptional project delivery - TCS
- **2X On the Spot Award** for time-sensitive deliverables - TCS