

NIKHIL REDDYVARI

DATA ENGINEER

Bangalore, India | nikhil.reddy0005@gmail.com | +91 8688127709

CAREER SUMMARY

Data Engineer with 4.5 years of experience building production-grade data platforms for Fortune 500 clients, including GlaxoSmithKline. Hands-on expertise with Azure Databricks, ADF, Synapse, PySpark, Delta Lake, DBT, and Power BI. Currently owning end-to-end streaming pipelines processing 5,000+ events/min, Medallion Lakehouse architectures, and UC governance at scale. Proven ability to bring enterprise-grade data engineering to high-impact product teams.

TECHNICAL SKILLS

Core Stack: Azure Databricks, Azure Data Factory (ADF), Azure Synapse Analytics, ADLS Gen2, Delta Lake, DBT

Data Engineering: Apache Spark, PySpark, Spark SQL, Structured Streaming, Medallion Architecture, Unity Catalog, API Integration, SFTP

Languages & Modeling: Python, SQL, Star & Snowflake Schema, SCD Type 2, Fact & Dimension Models, OLAP/OLTP

DevOps & Reporting: Azure DevOps, Git, CI/CD, Power BI, Data Quality Frameworks, Data Lineage, Access Control

WORK EXPERIENCE

Software Engineer (Data Engineer) | Tech Mahindra – GSK Client 21 March 2025 – Present

- Designed and owned end-to-end real-time streaming architecture processing **5,000+ occupancy sensor events/min across GSK's global offices**. This feeds the GOMA Power BI dashboard used by senior leadership for global workspace utilization analytics and capacity planning.
- Designed an Azure Data Factory pipeline to migrate **1.3 million records from a ServiceNow API** to a SQL Database, ensuring seamless downstream accessibility and high-volume data availability.
- Engineered a centralized enterprise mapping framework unifying many data sources (Snow, Serraview, Badge, WiFi, HR, AppSpace, KODE, XY) utilizing both **REST API and secure SFTP ingestion** patterns.
- Designed and implemented extended **Medallion Architecture** (Raw/Foundation/Trusted/Unified) on Delta Lake with **hash-based SCD Type 2**. Enabled full historical auditability, lineage, and change detection across critical enterprise datasets.
- Led the Hive Metastore to **Unity Catalog migration for all Databricks** workloads. Strengthened data governance, RBAC, and cross-workspace lineage for a highly regulated environment.
- Built **automated real-time pipelines for license assignment and CCURE access** provisioning. Implemented enterprise-grade monitoring, exception handling, and data quality validation frameworks.

Data Engineer I | Safran Engineering Services 03 Feb 2022 – 02 Jan 2025

- Architected and delivered end-to-end Azure data pipelines (**ADF, Databricks, Synapse, ADLS Gen2**) for flight logs, engineering records, and sensor telemetry, orchestrating file transfers and automated data loading. Optimized PySpark ETL frameworks, achieving a 35% processing efficiency improvement.
- Designed Star/Snowflake dimensional models and Delta Lake Lakehouse architecture for 5TB+ aerospace datasets. Reduced storage costs by 18% via partitioning, Z-ordering, and compaction.
- Migrated legacy on-premises data warehouse to Azure. Improved overall data quality and reliability by 15% through reconciliation, root cause analysis, and automated validation frameworks.

EDUCATION

Bachelors in Technology (Aurora's Technological & Research Institute, Hyderabad) July 2018 – June 2021